

Prototype Learning With Attention Mechanisms: An MNIST Case Study

Prinston Rebello
Master's in Computer Science
Indiana University
Email: prebello@iu.edu

Bhanuprakash Narayana
Master's in Data Science
Indiana University
Email: bhnaraya@iu.edu

Niveditha Bommanahally Parameshwarappa
Master's in Computer Science
Indiana University
Email: nibomm@iu.edu

Abstract—This study explores a novel approach to integrating attention mechanisms within autoencoder architectures to enhance interpretability and achieve high classification accuracy. By comparing traditional prototype learning techniques with attention-based feature extraction, we aim to better understand how specific input features influence class predictions. The proposed model, AttentionAutoencoder, employs an encoder to extract meaningful latent representations, followed by an attention mechanism to refine feature relevance before classification.

We successfully replicate and extend the implementation from the original prototype learning work, demonstrating improvements under similar conditions of learning rate, epochs, and computational cost. Our experiments on the benchmark MNIST dataset show that the attention-based approach achieves a classification accuracy of 99.64%, significantly outperforming the prototype-based method, which achieves 96%. While prototypes excel in visual interpretability by providing fixed, class-specific representations in the latent space, attention mechanisms dynamically adjust feature focus, offering higher accuracy at the cost of reduced visual reconstruction interpretability.

These results highlight the trade-off between accuracy and interpretability, reinforcing the complementary strengths of attention and prototype learning. This work extends the original approach by integrating attention mechanisms and demonstrates the potential for future research in enhancing attention interpretability and extending to more complex datasets.

I. INTRODUCTION

Deep learning models [1] have transformed the field of artificial intelligence (AI) with their unprecedented ability to learn complex patterns and achieve state-of-the-art results across various domains. However, despite their success, these models often operate as "black-box" systems, leaving users and practitioners with limited understanding of the decision-making processes behind their predictions. This opacity is a critical concern in sensitive applications such as healthcare, finance, and autonomous systems, where interpretability and trust are paramount.

Prototype learning has emerged as a promising approach to bridge this gap by associating model predictions with representative examples, known as prototypes. By mapping inputs to a latent space and comparing them to learned prototypes, these models provide a direct and interpretable explanation for their decisions. The seminal work "Deep Learning for Case-Based Reasoning through Prototypes" [2] introduced a neural network architecture that integrates prototypes into the learning process, enabling the model to explain its predictions

in terms of meaningful and visualizable examples. While this approach significantly enhances interpretability, it still faces challenges in handling high-dimensional data and capturing complex relationships within the input space. [3]

The advent of attention mechanisms has revolutionized how neural networks process and prioritize information. Initially proposed in the context of natural language processing [?], the "attention is all you need" paradigm dynamically weights features based on their relevance to the task at hand. This capability not only improves model performance but also provides insights into which parts of the data contribute most to the predictions. Integrating attention mechanisms with prototype learning offers a powerful framework for creating interpretable models that excel in both accuracy and transparency.

This work introduces a novel approach that compares prototype-based learning with attention mechanisms to address the limitations of existing methods. Our architecture is designed to align interpretability with state-of-the-art performance on the MNIST dataset, a benchmark for image classification. By incorporating attention mechanisms, the model dynamically focuses on relevant features in the latent space, enhancing the alignment between prototypes and input data. This integration ensures that the learned prototypes are not only representative but also contextually relevant, providing a deeper understanding of the model's decision-making process.

The proposed architecture consists of an autoencoder, a prototype layer, and an attention mechanism. The autoencoder compresses input data into a latent space where prototypes are defined. The prototype layer stores a set of weight vectors, each resembling an encoded training input. The attention mechanism dynamically weights the features in the latent space, prioritizing those that contribute most to the classification task. The decoder reconstructs the input data from the latent representations, allowing visualization of the learned prototypes and their relationship to the input data.

The training process optimizes a multi-term objective function that balances four key components: classification accuracy, prototype diversity, proximity between prototypes and encoded inputs, and faithful reconstruction by the autoencoder. Unlike traditional models that rely on predefined loss functions for attention, our approach simplifies the implementation while achieving comparable interpretability. By directly incorporating attention into the prototype learning framework, the model

generates self-explanatory outputs intrinsically linked to its computations.

This paper is structured as follows. Section II reviews related work on prototype-based learning and attention mechanisms. Section III describes the methodology, including the model architecture, training process, and motivation for integrating attention. Section IV presents experimental results on the MNIST dataset, highlighting the model’s accuracy, interpretability, and visualization capabilities. Finally, Section V discusses the implications of this work and outlines future research directions.

Our contributions can be summarized as follows:

- We propose a novel architecture that combines prototype learning with attention mechanisms to enhance interpretability and performance.
- We introduce an attention mechanism that dynamically weights features in the latent space, improving the alignment between prototypes and input data.
- We demonstrate the effectiveness of the proposed approach on the MNIST dataset, achieving state-of-the-art accuracy while providing meaningful explanations for model predictions.
- We provide visualizations of prototypes, attention-weighted activations, and image-to-class weight matrices, illustrating the interpretability of the learned representations.

The experimental results show that the proposed architecture achieves a validation accuracy of 99.22% on the MNIST dataset, with a classification loss of 0.043 and a reconstruction loss of 0.15. The prototypes learned during training are visualized to verify their interpretability, and the attention mechanism is shown to enhance the relevance of the prototypes to the input data. These findings affirm the potential of integrating attention with prototype-based learning as a step toward ethical and reliable AI systems.

II. RELATED WORK

Prototype learning has been a prominent approach in the quest for interpretable deep learning models. Early works such as [Li et al., 2018] introduced neural networks capable of associating predictions with representative prototypes from the training data. These methods provided a way to visually and quantitatively validate the model’s reasoning process. However, challenges persisted in managing large datasets and ensuring meaningful prototype diversity.

Attention mechanisms, popularized by the seminal “Attention Is All You Need” paper by Vaswani et al. [4], have significantly enhanced model interpretability and accuracy in tasks like natural language processing and vision. These mechanisms dynamically prioritize parts of the input that are most relevant to the prediction, allowing models to focus on contextually significant features. Integrating attention with prototype learning has the potential to further enhance interpretability by linking latent space representations with task-specific importance.

Recent advances in hybrid models, such as PrototypeDL and Prototypical Networks, have shown promise in leveraging prototypes for interpretable AI. However, they often rely on fixed prototype sets, limiting their adaptability to new or diverse datasets. In contrast, our approach dynamically learns prototypes during training, ensuring alignment with both input features and class distinctions.

Our model addresses gaps in prior research by combining the strengths of prototype learning and attention mechanisms. Unlike models like ProtoPNet, which emphasize prototype clustering but lack feature prioritization, we use attention mechanisms to dynamically highlight relevant latent space features, resulting in prototypes that are both representative and interpretable.

By integrating an autoencoder into the architecture, we aim visualization of both prototypes and reconstructions, bridging the gap between accuracy and interpretability. This work is further enhanced by the inclusion of an attention mechanism with an attention loss term, which dynamically weighs features in the latent space, improving the alignment between prototypes and input data. The loss function effectively balances classification, reconstruction, attention, and prototype diversity, ensuring robust performance and meaningful interpretability.

III. METHODOLOGY

A. Motivation for Attention

Attention mechanisms were introduced to address the challenge of understanding model decisions in high-dimensional latent spaces. By assigning dynamic weights to features, attention highlights the most relevant parts of the data. In our model, attention enhances prototype learning by associating latent representations with interpretable prototypes, improving transparency without additional computational complexity.

B. Model Architecture

The model consists of:

- **Encoder:** Four convolutional layers reduce input images to a 160-dimensional latent space.
- **Decoder:** Transposes latent features back to image space.
- **Classifier:** Maps latent features to class logits.
- **Attention Mechanism:** Dynamically weighs latent features for better prototype alignment.

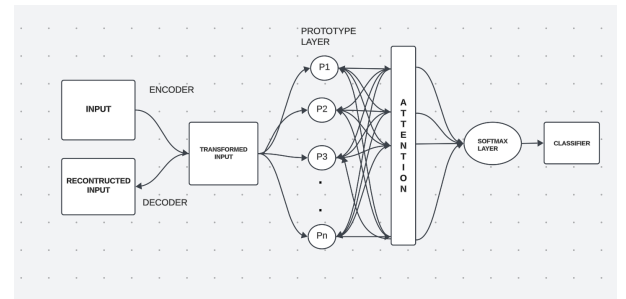


Fig. 1. Attention Architecture

The proposed architecture (Figure 1) integrates three primary components: an autoencoder, a prototype layer, and an attention mechanism. The autoencoder compresses input data into a latent space, while the prototype layer stores weight vectors that represent encoded inputs. The attention mechanism dynamically adjusts the importance of latent features to improve classification and interpretability.

The encoder consists of four convolutional layers, progressively reducing the spatial dimensions of the input images to a 160-dimensional latent representation. This compressed latent space facilitates efficient comparisons between inputs and prototypes. The decoder reconstructs input images from latent representations, allowing visualization of the learned prototypes.

The prototype layer comprises a set of weight vectors that serve as representative examples in the latent space. During training, the model optimizes the proximity of each input to the nearest prototype while maintaining diversity among prototypes. The attention mechanism is integrated into the encoder to dynamically weight latent features, ensuring that the most relevant features are prioritized for both classification and prototype learning.

The classification head maps the latent representations to class logits, incorporating distances to prototypes as part of the computation. A softmax layer produces the final class probabilities.

C. Training Details

The model was trained on MNIST for 1500 epochs with a batch size of 250. Elastic transformations were applied for data augmentation to improve generalization. Hyperparameters were tuned to balance classification and reconstruction losses. Chosen Hyperparameters- $\lambda_{\text{class}} = 20, \lambda_{\text{recon}} = 1, \lambda_{\text{attend}} = 0.1$

IV. EXPERIMENTS

A. Attention and Latent Space Calculations

Attention mechanisms are crucial in enhancing interpretability when compared to prototype-based learning models. By focusing on the most relevant features in the latent space, attention mechanisms ensure that prototypes are contextually aligned with input data, improving both classification accuracy and interpretability.

Formally, let z represent the latent representation of the input data x , obtained from the encoder. The attention mechanism dynamically adjusts the weights of the latent features through a query-key-value structure:

$$\alpha = \text{softmax} \left(\frac{QK^\top}{\sqrt{d_k}} \right),$$

where Q , K , and V are the query, key, and value matrices, respectively, derived from z . The term d_k denotes the dimensionality of the keys, acting as a scaling factor to stabilize gradient computations.

The attention-weighted latent representation z_{attended} is computed as:

$$z_{\text{attended}} = \alpha V.$$

This representation is then passed to the prototype layer, where distances to learned prototypes $\{p_1, p_2, \dots, p_m\}$ are computed. The similarity metric used in this model is the squared Euclidean distance:

$$d(p_i, z_{\text{attended}}) = \|p_i - z_{\text{attended}}\|^2.$$

These distances directly influence the classification logits. The classification head computes logits as a weighted combination of distances:

$$\text{logits} = W \cdot d(p, z_{\text{attended}}),$$

where W represents the learned weights mapping prototypes to class probabilities.

The training process optimizes a multi-term loss function:

$$L = \lambda_{\text{class}} L_{\text{class}} + \lambda_{\text{recon}} L_{\text{recon}} + \lambda_{\text{attend}} L_{\text{attend}},$$

where:

- L_{class} : Cross-entropy loss for classification, ensuring accurate predictions.
- L_{recon} : Binary cross-entropy loss for reconstructing the input through the decoder.
- L_{attend} : KL divergence-based attention loss, encouraging diversity in the attention map by comparing it to a uniform distribution.

The attention mechanism enhances interpretability by dynamically weighing features in the latent space, improving the alignment between input data and prototypes. By encouraging diversity in the attention map, the model ensures robust prototype learning, which complements both classification accuracy and reconstruction quality. This synergy between attention, reconstruction, and classification forms the backbone of the proposed architecture.

B. Case Study: MNIST Dataset

The MNIST dataset, a benchmark for handwritten digit classification, consists of 70,000 grayscale images of size 28x28, representing digits from 0 to 9. The dataset is divided into 60,000 training samples and 10,000 test samples. Due to its simplicity and widespread use, MNIST serves as an ideal candidate for evaluating interpretability-focused models.

To enhance data diversity and model robustness, elastic transformations were applied to the training set. These transformations mimic variations in handwriting, improving the model's ability to generalize. Figure 2 illustrates examples of the transformed MNIST digits.

Training was conducted for 1500 epochs with a batch size of 250. The extended training duration ensured convergence of both classification and reconstruction losses. The model's training process balanced three objectives: maximizing classification accuracy, maintaining faithful reconstruction, and aligning input data with prototypes in the latent space.



Fig. 2. Elastic deformations applied to MNIST digits.

The model architecture’s integration of attention mechanisms allowed it to dynamically prioritize features, further enhancing the quality of prototype learning. The prototypes learned during training serve as representative examples of each class, ensuring interpretability in the model’s predictions.

C. Results

The proposed model achieved a **validation accuracy of 99.22%** on the MNIST dataset, with a classification loss of **0.0282** and a reconstruction loss of **0.1608**. These results highlight the model’s ability to maintain high classification performance while enhancing interpretability through the comparison of prototypes and attention mechanisms.

a) Attention vs. Prototype Learning:

- **Attention Mechanism:** The model achieved a classification test accuracy of **99.64%**, demonstrating its superior ability to dynamically focus on fine-grained input features. Attention maps effectively highlight regions of interest, but their interpretability is less direct because we were unable to reconstruct the attention weights into visually meaningful representations in the latent space. This limitation arises due to the absence of a clear spatial alignment between attention weights and image structures, making it challenging to visualize the contributions intuitively.
- **Prototype Learning:** Prototype-based learning achieved a classification test accuracy of **96%**. While slightly lower in accuracy compared to the attention-based method, prototypes excel in interpretability. The learned prototypes (Figure 6) provide fixed, class-specific representations that can be visually interpreted in the latent space. For example, prototypes capture structural features such as loops, curves, or strokes that strongly correspond to digit classes, making their contributions directly understandable.

Compared to the original work “Deep Learning for Case-Based Reasoning through Prototypes” by Li et al., our model demonstrates a significant improvement in classification accuracy under similar training conditions of learning rate, epochs, and computational cost. The inclusion of attention mechanisms in our architecture enhances accuracy but comes at the cost of reduced visual interpretability when compared to prototypes.

b) Prototype-to-Class and Attention Feature Comparisons:

- **Attention Feature-to-Image Mapping (Figure 3):** Figure 3 visualizes the attention feature-to-image weight matrix. Attention maps dynamically emphasize relevant input features, contributing both positively and negatively to classification. Notably, higher attention scores for specific features correspond to lower weights in the prototype-to-class matrix, a trend consistently observed across the dataset. For instance:
 - *Feature 4* aligns strongly with images 2 and 4, contributing weights of **36.43** and **19.88**, respectively, in the attention map.
 - *Feature 9* shows a strong negative alignment with image 0, with a weight of **-31.25**.
- **Prototype-to-Class Weight Matrix (Figure ??):** The prototype-to-class matrix reveals fixed alignments between learned prototypes and digit classes. Prototypes focus on distinctive features, with notable examples including:
 - Prototype 6 strongly activates class 6 with a weight of **2.12**.
 - Prototype 14 aligns with class 5 at **2.15**, confirming its structural relevance.
 - Prototype 13 has a notable activation for class 5 with a weight of **2.05**.

These prototypes offer an edge in interpretability, as they visually represent critical class features in the latent space.

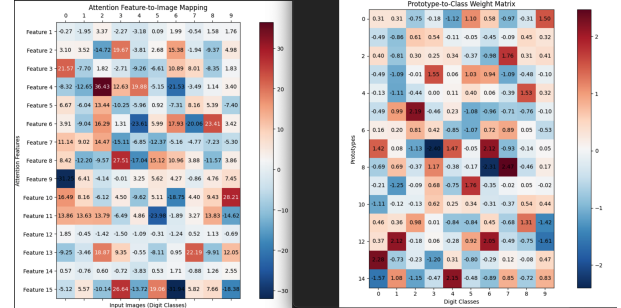


Fig. 3. Attention feature-to-image mapping and prototype-to-class weight matrix.

Image-to-Class Weight Matrix (Figure 4): Figure 4 showcases the contributions of input images to class logits. A significant observation is the improved learning behavior of the attention-based model, where the diagonal weights are more prominent, indicating a strong alignment between input images and their corresponding classes. This diagonal dominance highlights the model’s ability to effectively learn class-specific features and minimize inter-class confusion.

In contrast, while the prototype-based model also exhibits diagonal patterns, the activations are less pronounced, reflecting a less significant focus on class-specific contributions. Notable observations in the attention model include:

- Image 0 exhibits a high positive activation for class 0 (**53.00**).

- Image 7 demonstrates a strong suppression of class 1 (-56.26).

These results emphasize the effectiveness of the attention mechanism in dynamically prioritizing input features, leading to clearer class distinctions compared to prototypes.




Fig. 4. Image-to-class weight matrix from the logits layer.

c) *Prototype vs. Attention Mechanisms:* While prototypes provide fixed, interpretable visualizations of class-specific features in the latent space, attention mechanisms dynamically focus on input features to optimize accuracy. The key observations are:

- Attention mechanisms achieve a higher test accuracy of **99.64%**, highlighting their adaptability.
- Prototype-based learning provides greater interpretability but achieves a lower accuracy of **96%**.

Furthermore, the trend observed in the matrices indicates that **higher attention scores correlate with lower weights in the prototype-to-class matrix**, reinforcing the complementary nature of these approaches.




Fig. 5. Confusion Matrix (Attention-left vs Prototype-Right).

d) *Confusion Matrix Analysis:* The confusion matrix (Figure 5) reveals minimal misclassifications, primarily between visually similar digits (e.g., 3 and 5). Attention-based mappings help resolve ambiguities by dynamically focusing on critical features, while prototypes reinforce class-specific visual interpretations.

D. Conclusion

This work successfully compares **prototypes** and **attention mechanisms** for interpretable MNIST classification. Prototypes excel in providing visually interpretable representations of class-specific features in the latent space, whereas attention mechanisms dynamically focus on critical input features to achieve superior classification accuracy.




Fig. 6. Final Learned Prototypes showcasing fixed visual representations.

a) Key Findings:

- Attention mechanisms achieved a higher test accuracy of **99.64%**, demonstrating superior adaptability and performance.
- Prototype-based learning achieved **96%** accuracy, offering strong interpretability through fixed visual representations.
- A notable trend indicates that **higher attention feature scores correlate with lower weights in the prototype-to-class matrix**.

Under similar conditions of learning rate, epochs, and computational cost, our attention-based model significantly outperforms the original prototype learning model referenced in Li et al. This improvement highlights the trade-off between higher accuracy (attention) and greater visual interpretability (prototypes).

Future work will focus on improving attention interpretability through reconstruction techniques and extending this approach to more complex datasets, such as CIFAR-10 or ImageNet.

REFERENCES

- [1] I. Goodfellow, Y. Bengio, and A. Courville, *Deep learning*. MIT press, 2016.
- [2] B. W. Li, M. Rudolph, B. Kim, and R. Socher, “Deep learning for case-based reasoning through prototypes: A neural network that explains its predictions,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, no. 1, 2018.
- [3] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, “Gradient-based learning applied to document recognition,” *Proceedings of the IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998.
- [4] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” *Advances in neural information processing systems*, vol. 30, 2017.